

# Nevronska mreža

Tomaž Tomažič (63100281)

17. maj 2015

## 1 Uvod

Cilj naloge je bil cim bolj pravilno napovedati eno izmed devetih skupin, kateri pripada nek izdelek.

## 2 Podatki

Podatki so imeli 93 značilk katere so bile anonime. Vsi podatki so zasegali diskretne vrednosti in so bili nenegativni. Na voljo je bilo 50000 učnih primerov, v katerih sta močno prevladala drugi in šesti razred.

## 3 Metode

Najprej sem napisal kodo za nevronske mreže po članku Šparse autoencoder; ki je delovala samo za 2 skriti plasti. Ko sem se prepričal, da ta koda deluje in da z numeričnim preverjanjem odvod skoraj enake rezultate kot analitični, sem napisal verzijo, ki deluje za poljubno število skritih plasti.

Na zadnjem nivoju sem zatem zamenjal aktivacijsko funkcijo za softmax funkcijo, saj želim imeti na koncu verjetnosti za določen razred, ki se seštevajo v 1. Poleg nove aktivacije na zadnjem nivoju sem moral zamenjati cenilno funkcijo za cenilko softmax regresije in posodobitev napake na zadnjem nivoju v odvodu.

Implementiral sem tudi dropout tehniko, da sem zmanjšal overfit nevronske mreže. Nevrone izlocim iz mreže po s predlagano verjetnostjo po bernoullijevi distribuciji. Na vhodnem nivoju odstranim nevrone z vrjetnostjo 0.2 na ostalih nivojih pa to verjetnost postopoma povečujem. Nevroni se naključno izberejo ob vsakem klicu backprop funkcije. Za racunanju cenilne funkcije se pa uporabijo vsi nevroni.

Za vsako tehniko izboljšav sem preveril pravilnost gradienta z numeričnim odvajanjem (funkcija `test_grad()`) in preveril model s precnim preverjanjem. To sem naredil taok da sem podatke razdelil na 2 dela. En del je ostal vedno neviden za končno evalvacijo modela. Za izbiro najboljše lambde sem uporabil funkcijo iz paketa `sklearn - cross_val_score()` v kateri se je testiralo podatke v treh delih/foldih.

Ker se je pri softmax domaci nalogi podatke splačalo normalizirati, sem to tudi tu poiskusil, vendar sem vedno dobil veliko slabše rezultate in sem zato to tehniko opustil.

## 4 Rezultati

Ker je v nevronskih mrežah ogromno parametrov, katere je potrebno izbrati, je optimalna izbira praktično časovno nemogoča. Zato sem se odločil, da vse modele za primerjavo testiram pri enaki arhitekturi (93, 20, 20, 9) z dvema skritima nivojema in poiščem čim boljši regularizacijski parameter  $\lambda$ . Rezultati modelov po prečnem preverjanju so podani v tabeli 1. Zraven so, za enake modele, rezultati iz spletnega strežnika.

Tabela 1: Rezultati uporabljenih tehnik v nevronski mreži.

tehnika	prečno preverjanje	oddaja	lambda
sparse autoencoder	0.615234227501	0.61665	0.000001
softmax	0.574657559294	0.58094	0.00027
dropout	0.560926567503	0.56929	0.0005

Na strežniku sem poizkušal nekaj oddaj tudi z drugačnimi arhitekturami, vendar rezultati niso bili boljši.

## 5 Izjava o izdelavi domače naloge

Domačo nalogo in pripadajoče programe sem izdelal sam.

# Priloge

## A Programska koda

Programska koda je v 4 datotekah, zato bom v komentar napisal na začetek ime datoteke

```
1 #nn.py
2 from NeuralNet import NeuralNet
3 import IO
4 import CV
5 import Orange
6 import numpy as np
7
8 # layers = (4, 10, 10, 3)
9 # iris = Orange.data.Table("iris")
10 # CV.simple_prediction_test(iris.X, iris.Y, layers)
11 # CV.eval_model(iris.X, iris.Y, layers)
12
13 layers = (93, 20, 20, 9)
14 # X, Y = IO.readFile()
15 evalX = IO.readTestFile()
```

```

X, Y = IO.readFile('train.csv')
17
# CV.eval_model(X, Y, layers)
19 CV.predict(X, Y, evalX, layers, "result-softmax-20-20")

21
#CV.py
23 from NeuralNet import NeuralNet
import IO
25 import numpy as np
import Orange
27 from sklearn import preprocessing, metrics, grid_search
from sklearn.cross_validation import cross_val_score, ShuffleSplit,
train_test_split
29
def find_lambda(X, Y, nnl):
31     '''get best lambda'''

33     lambdas = list(map(float, np.linspace(0.0000001, .001, 5))) # +\
# list(map(float, np.linspace(.1, 2, 5))) +\
35     # list(map(float, np.linspace(2, 100, 10)))

37
    cv_split = ShuffleSplit(Y.shape[0], n_iter=3, test_size=0.3, random_state
=42)
39     cv_score = lambda l: -np.mean(cross_val_score(nnl, X, Y, \
cv=cv_split, fit_params={'lambda_': l}, scoring = 'log_loss', n_jobs=3)
41     )

43     best_lambda = min([(cv_score(l), l) for l in lambdas])
print("best labmda", best_lambda[1], " got mean score ", best_lambda[0], "
for 3 folds")
45     return best_lambda[1]

47 def eval_model(X, Y, layers, lambda_=None):
'''evaluate model with best lambda on unseen data'''
49
    X_train, X_test, Y_train, Y_test = train_test_split(
51     X, Y, test_size=0.2, random_state=42
    )

53
    nnl = NeuralNet(layers)
55     lambda_ = find_lambda(X_train, Y_train, nnl) if lambda_ is None else
lambda_
    nnl.fit(X_train, Y_train, lambda_=lambda_)
57     y = nnl.predict(X_test)
    result = metrics.log_loss(Y_test, y)
59     print("this model got score ", result, " with lambda ", lambda_, " and arch
", layers)
    return lambda_

61
def predict(X_train, Y_train, X, layers, filename="result"):
63     nnl = NeuralNet(layers)
    lambda_ = eval_model(X_train, Y_train, layers)
65     nnl.fit(X_train, Y_train, lambda_=lambda_)
    prediction = nnl.predict(X)
67     IO.savePrediction(prediction, filename)
    print("prediction finished")

69
def simple_prediction_test(X, Y, layers, lambda_=0.001):

```

```

71     nnl = NeuralNet(layers)
       X_train, X_test, Y_train, Y_test = train_test_split(
73         X, Y, test_size=0.2, random_state=42
       )
75     nnl.fit(X_train, Y_train, lambda_=lambda_)
       prediction = nnl.predict(X_test)
77     print(prediction)
       print(metrics.log_loss(Y_test, prediction))
79
# IO.py
81 import csv
import numpy as np
83 from sklearn import preprocessing

85 def readFile(path="data4_reduced.csv"):
       with open(path, newline='') as csvfile:
87         data = csv.reader(csvfile, delimiter=',')
         data = [i[1:-1] + [i[-1][-1]] for i in data]
89
       data.pop(0)
91       data = np.array(data).astype(int)
       dataX = data[:, :-1]
93       dataY = data[:, -1]-1
       return dataX, dataY
95
def readTestFile(path="test.csv"):
97     with open(path, newline='') as csvfile:
         data = csv.reader(csvfile, delimiter=',')
99         data = [i[1:] for i in data]

101     data.pop(0)
       data = np.array(data).astype(int)
103     return data

105 def savePrediction(p, name="result"):
       f = open(name + ".csv", 'w')
107       f.write("id,Class_1,Class_2,Class_3,Class_4,Class_5,Class_6,Class_7,Class_8
,Class_9\n")
       np.set_printoptions(suppress=True)
109       np.set_printoptions(precision=7)
       for i, j in enumerate(p):
111         f.write(str(i+1) + "," + ",".join(j.astype(str)) + "\n")
       f.close()
113
115
# NeuralNet.py
117
import numpy as np
119 import Orange
from scipy import optimize, stats

121
# epsilon = 0.00000001
123 np.random.seed(42)

125 class NeuralNet:
       def __init__(self, arch):
127         self.arch = arch
         self.theta_shape = np.array([(arch[i]+1, arch[i+1])
129           for i in range(len(arch)-1)])

```

```

131         ind = np.array([s1*s2 for s1, s2 in self.theta_shape])
132         self.theta_ind = np.cumsum(ind[:-1])
133         self.theta_len = sum(ind)
134
135     def init_thetas(self, epsilon=1):
136         return np.random.rand(self.theta_len) * 2 * epsilon - epsilon
137
138     def shape_thetas(self, thetas):
139         t = np.split(thetas, self.theta_ind)
140         return [t[i].reshape(shape) for i, shape in enumerate(self.theta_shape)]
141
142     def h(self, a, thetas):
143         """feed forward, prediction"""
144         thetas = self.shape_thetas(thetas)
145         for theta in thetas[:-1]:
146             a = self.g(self.add_ones(a).dot(theta))
147
148         # without softmax
149         # return self.g(self.add_ones(a).dot(thetas[-1]))
150
151         # softmax last layer
152         s = self.softmax(self.add_ones(a), thetas[-1])
153         return s
154
155     def J(self, thetas):
156         Y = self.Y
157         theta = self.shape_thetas(thetas)
158         h0 = np.maximum(np.minimum(self.h(self.X, thetas), 1 - 1e-15), 1e-15)
159
160         # logistic output layer
161         # J = .5 * np.sum(np.power(h0 - Y, 2)) / self.m
162         # reg = np.sum([np.sum(t[1:].dot(t[1:].T)) for t in theta])
163         # return J + self.lambda_ / 2.0 * reg
164
165         # cost for softmax
166         J = -np.sum(Y*np.log(h0))
167         # remove bias in theta for regularization
168         reg = np.sum([np.sum(np.power(t[1:], 2)) for t in theta])
169         return (J + self.lambda_ / 2.0 * reg) / self.m
170
171     def grad_approx(self, thetas, e=1e-2):
172         return np.array([(self.J(thetas+eps) - self.J(thetas-eps))/(2*e)
173                         for eps in np.identity(len(thetas)) * e])
174
175     def softmax(self, activation, theta):
176         z_last = activation.dot(theta)
177         s = np.exp(z_last - np.max(z_last, axis=1)[:,None])
178         s /= np.sum(s, axis=1)[:, None]
179         return s
180
181     def backprop(self, thetas):
182         X = self.X
183         Y = self.Y
184         theta = self.shape_thetas(thetas)
185         dropout = self.dropout_init()
186
187         #####
188         # feed forward with history
189         #####

```

```

189     # act = [self.add_ones(X)] #activation on first layer (
a1)
191     # for t in theta[:-1]: #activation on middle layers
        # act.append( self.add_ones(self.g(act[-1].dot(t))) )

193     #with dropout
act = [self.add_ones(X) * dropout[0]] #activation on first layer (a1)
195     # act = [self.add_ones(X)] #activation on first layer (a1)
for i, t in enumerate(theta[:-1]): #activation on middle layers
197         act.append( self.add_ones(self.g(act[-1].dot(t))) * dropout[i+1] )
        # act.append( self.add_ones(self.g(act[-1].dot(t))) )
199
201     # logistic output if u have no softmax
        # act.append( self.g(act[-1].dot(theta[-1])) ) #activation on last
layer

203     #softmax for last layer
s = self.softmax(act[-1], theta[-1])
205     act.append(s)

207     #reverse for backprop
act.reverse(), theta.reverse(), dropout.reverse()
209
211     #####
# backpropagation
213     #####
d = [(act[0]-Y)] #error for softmax
        # d = [-(Y-act[0]) * act[0] * (1-act[0])] #error on normal
last layer
215     for a, t in zip(act[1:-1], theta[:-1]): #errors on middle
layers
        d.append((t.dot(d[-1].T).T * (a * (1-a)))[:, 1:])

217
219     #without regularization
        # new theta for every layer reversed and flatten back to vector
        # D = [ a.T.dot(e).ravel() for a, e in zip(reversed(act[1:]), reversed(
d))]
221     # return np.hstack(D) / self.m

223     D = [(a * drop).T.dot(e) / self.m for a, e, drop in zip(act[1:], d,
dropout)]
        # D = [a.T.dot(e) / self.m for a, e in zip(act[1:], d)]
225
227     # with regularization
        _l = self.lambda_
        for i in range(len(theta)): theta[i][0] = 0 # remove bias
229     D = [ (d + _l * t).ravel() for d, t in zip(reversed(D), reversed(theta)
)]

231     return np.hstack(D)

233     #####
# step by step code for 2 hidden layers only.
# usefull to debug upper code
235     # here is without regularization, softmax and dropout
#####
237     # a1 = self.add_ones(X)

239     # z2 = a1.dot(theta[0])
        # a2 = self.add_ones(self.g(z2))
241

```

```

243     # z3 = a2.dot(theta[1])
244     # a3 = self.add_ones(self.g(z3))
245
246     # z4 = a3.dot(theta[2])
247     # a4 = self.g(z4)
248
249     # d4 = (a4-Y)
250     # d3 = (theta[2].dot(d4.T).T * (a3 * (1-a3)))[:, 1:]
251     # d2 = (theta[1].dot(d3.T).T * (a2 * (1-a2)))[:, 1:]
252
253     # T3 = a3.T.dot(d4) / self.m
254     # T2 = a2.T.dot(d3) / self.m
255     # T1 = a1.T.dot(d2) / self.m
256
257     # return np.hstack((T1.flat, T2.flat, T3.flat))
258
259 def fit(self, X, y, **kwargs):
260     # init variables
261     self.lambda_ = kwargs['lambda_'] if 'lambda_' in kwargs else 0
262     self.X, self.y = X, y
263     self.Y = np.eye(self.arch[-1])[self.y.astype(int)]
264     self.m = self.X.shape[0]
265     thetas = self.init_thetas()
266
267     # self.test_grad(thetas)
268
269     #get new thetas
270     thetas, fmin, d = optimize.fmin_l_bfgs_b(func=self.J, x0=thetas, fprime
=
self.backprop)
271     self.thetas = thetas
272     if d['warnflag'] != 0: print("ERROR!!, LBFGS is not working")
273     return self
274
275 def test_grad(self, thetas):
276     # print(self.init_thetas)
277     approx = np.absolute(self.grad_approx(thetas))
278     real = np.absolute(self.backprop(thetas))
279     print( np.absolute(sum((approx - real))))
280     print( np.sum((approx - real)**2))
281     # print( (np.absolute(approx - real) < 1e-4).all())
282
283 def dropout_init(self):
284     dropout = []
285     m = self.m
286     passtrough = .8
287     for i, a in enumerate(self.arch[:-1]):
288         dropout.append(stats.bernoulli.rvs(passtrough, size=(self.m, a+1)))
289         passtrough -= 0.05
290         dropout[-1][:,0] = 1
291         #disable dropout
292         # dropout[-1] = 1
293     return dropout
294
295 def g(self, z):
296     return 1/(1+np.exp(-z))
297
298 def add_ones(self, X):
299     return np.column_stack((np.ones(len(X)), X))
300
301 def get_params(self, deep = False):

```

```
301     '''used for CV'''
302     return {'arch': self.arch}
303
304     def predict_proba(self, X):
305         '''used for CV'''
306         return self.h(X, self.thetas)
307
308     def predict(self, X):
309         '''used for CV'''
310         return self.predict_proba(X)
```