Gradientni boosting

Tomazic Tomaz (63100281)

31. maj 2015

1 Uvod

Cilj naloge je bil napovedati eno izmed devetih skupin, kateri pripada nek izdelek.

2 Podatki

Podatki so imeli 93 značilk katere so bile anonime. Vsi podatki so zasegali diskretne vrednosti in so bili nenegativni. Na voljo je bilo 50000 učnih primerov, v katerih sta močno prevladala drugi in šesti razred.

3 Metode

Implementiral sem gradientni boosting ki za funkcijo izgube uporablja KL-divergenco (Kullback Leibler divergence). Razred GradBoost ob inicializaciji sprejme kot parameter šibak regresijski model, ki mora imeti implementirani funkciji 'predict' in 'fit'. Kot parameter sprejme tudi ime funkcije izgube. Trenutno je implementirana samo KL izguba. Razred 'GradBoost' deluje za vecrazredno klasifikacijo.

V fit metodi kot osnovno ucenje najprej izracunam frekvenco razredov nato pa izboljsujem napoved z regresijskim drevesom oz podanim ucnim modelom.

Napovedne modele sem ovrednotil z precnim preverjanjem v katerem sem iskal najboljse parametre za dane podatke. Ti parametri so: stevilo modelov za posamezni razred, globina regresijskega drevesa in regularizacijski parameter 'ro'. V precnem preverjanju sem poizkusal 30 modelov med 10 in 1000, globine drevesa od 1 do 5 in 5 regularizacijskij prametrov med 0.01 in 1.

Za zdruzevanje napovedi nisem uporabil metode stacking ampak samo povprecenje rezultatov mojih nevronskih mrez in rezultate gradient boostinga, ker sem vseeno pricakoval nekoliko boljse rezultate na lestvici.

4 Rezultati

Presenetilo me je da sem najboljse rezultate dobil z regularizacijskim parametrom vrednosti 1.

V tabeli 1 so najboljsi rezultati pri izbiri paramerov za doloceno globino drevesa. V tabeli-2 pa rezultati precnega preverjanja

Tabela 1: Rezultati povprecja treh foldov pri razlicnih parametrih.

st iteracij	globina drevesa	ro	rezultat
500	1	1	0.66513236171212353
500	2	1	0.57993880037618517
500	3	1	0.54497348832603998
500	4	1	0.53358386235816502
319	5	1	0.53697762802747351
625	4	1	0.53264532029139089

Tabela 2: Rezultati precnega preverjanja.

ime	st iteracij	globina drevesa	ro	vrednost modela	oddaja na strezniku (9b7)
boosting	625	4	1	0.51243599428	0.51648
nevronske mreze		51	0.003	0.54092656750	0.53926
povprecenje					0.49849

Rezultati ocitno kazejo da je boosting boljsa metoda v primerjavi z nevronskimi mrezami za dane podatke.

5 Izjava o izdelavi domače naloge

Domačo nalogo in pripadajoče programe sem izdelal sam.

Priloge

A Programska koda

```
import numpy as np
import sklearn
import copy
from Frequency import Frequency

class GradBoost:
    """Gradient Boosting for Classification."""

    def __init__(self, learner, n_estimators=100, loss="KL", epsilon=1e-5):
        self.n_estimators = n_estimators
        self.learner      = learner
        self.epsilon      = epsilon

```

```
        losses = {
            "huber": self.grad_huber_loss,
            "squared": self.grad_squared_loss,
            "abs": self.grad_abs_loss,
            "KL": self.kl_loss
        }
        self.loss = losses[loss]

    def grad_squared_loss(self, y, f):
        """Negative gradiant for squared loss."""
        return y - f

    def grad_abs_loss(self, y, f):
        """Negative gradient for absolute loss."""
        return np.sign(y - f)

    def grad_huber_loss(self, y, f, delta=0.5):
        """Negative gradient for Huber loss."""
        r0 = y - f
        r1 = delta * np.sign(r0)
        return np.vstack((r0, r1)).T[np.arange(y.shape[0]), (np.abs(r0)>delta).
    astype(int)]

    def kl_loss(self, y, f):
        """Negative gradient for kullback leibler."""
        return y - f

    def fit(self, X, Y, ro=1):
        num_c = Y.shape[1]

        #average
        models = [[Frequency().fit(X, Y[:, i]) for i in range(num_c)]]
        F      = np.hstack([i.predict(X) for i in models[0]])

        grad = self.loss(Y, F)

        for i in range(self.n_estimators):

            models.append([])

            for j in range(num_c):
                #fit regression tree
                dtc = copy.copy(self.learner)
                dtc.fit(X, grad[:, j])
                models[i+1].append(dtc)

                F[:, j] += dtc.predict(X) * ro

            P    = self.softmax(F)
            grad = self.loss(Y, P)

        self.models = np.array(models)
        return self


    def softmax(self, f):
        s = np.exp(f - np.max(f, axis=1)[:,None])
        s /= np.sum(s, axis=1)[:, None]
        return s
```

```
    def test_grad(self, Y):
        F = np.zeros(Y.shape)
        P = self.softmax(F)
        grad = self.loss(Y, P)

        # print(self.init_thetas)
        approx = np.absolute(self.grad_approx(Y, P))
        real = np.absolute(self.loss(Y, P))
        print( np.absolute(sum((approx - real))) )
        print( np.sum((approx - real)**2) )

    def KL(self, Q, P):
        return np.sum(P * np.log(P / (np.exp(Q)/np.sum(Q))))

    def grad_approx(self, Y, P, e=1e-4):
        num_grad = np.zeros_like(Y)
        perturb = np.zeros(Y.shape[0])
        for j in range(Y.shape[1]):
            for i in range(Y.shape[0]):
                perturb[i] = e
                j1 = self.KL(Y[:,j], P[:,j] + perturb)
                j2 = self.KL(Y[:,j], P[:,j] - perturb)
                num_grad[i, j] = (j1 - j2) / (2. * e)
                perturb[i] = 0
        return self.softmax(num_grad)

    def predict_proba(self, X):
        return self.predict(X)

    def predict(self, X):
        models = self.models
        num_c  = models.shape[1]

        P = np.hstack([i.predict(X) for i in models[0]])

        for i in range(num_c):
            P[:, i] += np.sum([j.predict(X) for j in models[1:, i]], axis=0)

        return self.softmax(P)

    def get_params(self, deep=False):
        return {
            "learner": self.learner,
            "n_estimators": self.n_estimators
        }


    ###################### Frequency.py ####################
    import numpy as np

    class Frequency:
        def fit(self, X, y):
            self.model = np.sum(y, axis=0) / y.shape[0]
            return self

        def predict(self, X):
            return np.vstack([self.model for i in range(X.shape[0])])

    ###################### CV.py ####################
```

```
133     from GradBoost import GradBoost
        import IO
135     import numpy as np
        import Orange
137     from sklearn import preprocessing, metrics, grid_search
        from sklearn.tree import DecisionTreeRegressor
139     from sklearn.cross_validation import cross_val_score, ShuffleSplit,
    train_test_split


141
        def cv_score(X, Y, split, n_estimators, depth):
143         learner = DecisionTreeRegressor(max_depth=depth)
            gb = GradBoost(learner, n_estimators)
145         return np.absolute(np.mean(cross_val_score(gb, X, Y, \
                cv=split, scoring='log_loss', n_jobs=4)
147         ))

149     def find_params(X, Y): #n_estimators && depth
            n_estimators = list(map(int, np.linspace(10, 500, 20)))
151         depths = list(map(int, np.linspace(1, 5, 5)))

153         cv_split = ShuffleSplit(Y.shape[0], n_iter=3, test_size=0.3,
    random_state=42)

155         scores = []
            for j in depths:
157             scores.append( min([(cv_score(X, Y, cv_split, i, j), i, j) for i in
     n_estimators]) )
                print(scores[-1])
159         print(scores)
            score = min(scores)
161         print("best n_estimators", score[1], "depth", score[2], " got mean
    score ", score[0], " for 3 folds")
            return (score[1], score[2])
163
        def eval_model(X, Y):
165         '''evaluate model with best lambda on unseen data'''

167         X_train, X_test, Y_train, Y_test = train_test_split(
                X, Y, test_size=0.2, random_state=42
169         )

171         n_estimators, depth = find_params(X_train, Y_train)
            learner = DecisionTreeRegressor(max_depth=depth)
173         gb = GradBoost(learner, n_estimators)
            gb.fit(X_train, Y_train)
175         y = gb.predict(X_test)
            result = metrics.log_loss(Y_test, y)
177         print("this model got score ", result, " with n ", n_estimators)
            return (n_estimators, depth)
179
        def predict(X_train, Y_train, X, filename="result"):
181         n_estimators, depth = eval_model(X_train, Y_train)

183         learner = DecisionTreeRegressor(max_depth=depth)
            gb = GradBoost(learner, n_estimators)
185         gb.fit(X_train, Y_train)
            prediction = gb.predict(X)
187         IO.savePrediction(prediction, filename)
            print("prediction finished")
```

```
189
    def simple_prediction_test(X, Y, lambda_=0):
191         learner = DecisionTreeRegressor(max_depth=4)
            gb = GradBoost(learner)
193
            X_train, X_test, Y_train, Y_test = train_test_split(
195             X, Y, test_size=0.2, random_state=42
            )
197         gb.fit(X_train, Y_train)
            prediction = gb.predict(X_test)
199         # print(prediction)
            print(metrics.log_loss(Y_test, prediction))
201
    ###################### IO.py ####################
203
    import csv
205 import numpy as np
    from sklearn import preprocessing
207
    def readFile(path="data4_reduced.csv"):
209         with open(path, newline='') as csvfile:
                data = csv.reader(csvfile, delimiter=',')
211             data = [i[1:-1] + [i[-1][-1]] for i in data]

213         data.pop(0)
            data = np.array(data).astype(int)
215         dataX = data[:, :-1]
            dataY = data[:, -1] #from 1 to 9
217         dataY = np.eye(np.max(dataY))[dataY.astype(int)-1]
            return dataX, dataY
219
    def readTestFile(path="test.csv"):
221         with open(path, newline='') as csvfile:
                data = csv.reader(csvfile, delimiter=',')
223             data = [i[1:] for i in data]

225         data.pop(0)
            data = np.array(data).astype(int)
227         return data

229 def savePrediction(p, name="result"):
            f = open(name + ".csv", 'w')
231         f.write("id,Class_1,Class_2,Class_3,Class_4,Class_5,Class_6,Class_7,
    Class_8,Class_9\n")
            np.set_printoptions(suppress=True)
233         np.set_printoptions(precision=7)
            for i, j in enumerate(p):
235             f.write(str(i+1) + "," + ",".join(j.astype(str)) + "\n")
            f.close()
237
    def normalize(X, X_test=None):
239         p = preprocessing.Normalizer().fit(X)
            X = p.transform(X)
241         if X_test is None:
                return (X)
243         else:
                X_test = p.transform(X_test)
245             return (X, X_test)


247
```

```
249     ###################### boost.py ####################

251     import Orange
        import numpy as np
253     import sklearn
        import CV
255     import IO
        from GradBoost import GradBoost
257     from sklearn.tree import DecisionTreeRegressor as DTC

259     numClasses = 3
        iris = Orange.data.Table("iris")
261     X = iris.X
        Y = np.eye(numClasses)[iris.Y.astype(int)]
263
        # print(GradBoost().fit(X, y))
265     # print(CV.simple_prediction_test(X, Y))
        # print(CV.eval_model(X, Y))
267


269     # X, Y = IO.readFile()
        # print(CV.simple_prediction_test(X, Y))
271     # print(CV.eval_model(X, Y))

273     X, Y = IO.readFile("train.csv")
        Y_test = IO.readTestFile()
275     CV.predict(X, Y, Y_test, "result")
```